

```
1 !pip install diffusers transformers accelerate peft safetensors ftfy pillow numpy matplotlib tqdm opencv-python scikit-image

Requirement already satisfied: diffusers in /usr/local/lib/python3.12/dist-packages (0.35.2)
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-packages (4.57.2)
Requirement already satisfied: accelerate in /usr/local/lib/python3.12/dist-packages (1.12.0)
Requirement already satisfied: peft in /usr/local/lib/python3.12/dist-packages (0.18.0)
Requirement already satisfied: safetensors in /usr/local/lib/python3.12/dist-packages (0.7.0)
Collecting ftfy
  Downloading ftfy-6.3.1-py3-none-any.whl.metadata (7.3 kB)
Requirement already satisfied: pillow in /usr/local/lib/python3.12/dist-packages (11.3.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (4.67.1)
Requirement already satisfied: opencv-python in /usr/local/lib/python3.12/dist-packages (4.12.0.88)
Requirement already satisfied: scikit-image in /usr/local/lib/python3.12/dist-packages (0.25.2)
Requirement already satisfied: importlib_metadata in /usr/local/lib/python3.12/dist-packages (from diffusers) (8.7.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from diffusers) (3.20.0)
Requirement already satisfied: huggingface-hub>=0.34.0 in /usr/local/lib/python3.12/dist-packages (from diffusers) (0.36.0)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.12/dist-packages (from diffusers) (2025.11.3)
Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-packages (from diffusers) (2.32.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (25.0)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-packages (from transformers) (6.0.3)
Requirement already satisfied: tokenizers<0.23.0,>=0.22.0 in /usr/local/lib/python3.12/dist-packages (from transformers) (0.2)
Requirement already satisfied: putil in /usr/local/lib/python3.12/dist-packages (from accelerate) (5.9.5)
Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.12/dist-packages (from accelerate) (2.9.0+cu126)
Requirement already satisfied: wcwidth in /usr/local/lib/python3.12/dist-packages (from ftfy) (0.2.14)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.5)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (2.9.0.post0)
Requirement already satisfied: scipy>=1.11.4 in /usr/local/lib/python3.12/dist-packages (from scikit-image) (1.16.3)
Requirement already satisfied: networkx>=3.0 in /usr/local/lib/python3.12/dist-packages (from scikit-image) (3.6)
Requirement already satisfied: imageio!=2.35.0,>=2.33 in /usr/local/lib/python3.12/dist-packages (from scikit-image) (2.37.2)
Requirement already satisfied: tifffile>=2022.8.12 in /usr/local/lib/python3.12/dist-packages (from scikit-image) (2025.10.16)
Requirement already satisfied: lazy-loader>=0.4 in /usr/local/lib/python3.12/dist-packages (from scikit-image) (0.4)
Requirement already satisfied: fsspec>=2023.5.0 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.34.0->diffu
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.34.0->diffu
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in /usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.34.0->diffu
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib) (1.1.5)
Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (75.2.0)
Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (1.14.6)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerate) (3.1.6)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cUBLAS-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cusparseelt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-nccl-cu12==2.27.5 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-nvshmem-cu12==3.3.20 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch>=2.0.0->accelerat
```

```
1 # imports and setup
2 import os
3 import random
4 import numpy as np
5 import cv2
6 import torch
7 import torch.nn.functional as F
8 from PIL import Image
9 from pathlib import Path
10 from tqdm import tqdm
11 from torch.utils.data import Dataset, DataLoader
12 from torchvision import transforms
13 from diffusers import StableDiffusionInpaintPipeline, DDPMscheduler
14 from peft import LoraConfig, get_peft_model
15 from torch.optim import AdamW
16 import matplotlib.pyplot as plt
17 from skimage.metrics import peak_signal_noise_ratio as psnr
18 from skimage.metrics import structural_similarity as ssim
19
```

```

20 # Configuration
21 IMAGE_SIZE = 512
22 DEVICE = "cuda" if torch.cuda.is_available() else "cpu"
23 print(f"Using device: {DEVICE}")
24
25 from google.colab import drive
26 drive.mount('/content/drive', force_remount=True)
27
28 # Paths
29 DAMAGED_DIR = Path("/content/drive/MyDrive/DL_Project/damaged images")
30 RAW_DIR = Path("/content/drive/MyDrive/DL_Project/Raw Images")
31 BASE_TRAIN_PATH = "/content/drive/MyDrive/DL_Project/train"
32 CLEAN_DIR = Path(BASE_TRAIN_PATH) / "clean images"
33 MASK_DIR = Path(BASE_TRAIN_PATH) / "mask"
34
35 # Create directories
36 CLEAN_DIR.mkdir(parents=True, exist_ok=True)
37 MASK_DIR.mkdir(parents=True, exist_ok=True)

```

Using device: cuda
Mounted at /content/drive

```

1 # preprocessing functions
2 def center_crop(img: Image.Image) -> Image.Image:
3     width, height = img.size
4     crop_size = min(width, height)
5     left = (width - crop_size) // 2
6     top = (height - crop_size) // 2
7     right = left + crop_size
8     bottom = top + crop_size
9     return img.crop((left, top, right, bottom))
10
11 def normalize_color(img: Image.Image) -> Image.Image:
12     img_cv = np.array(img)
13     img_yuv = cv2.cvtColor(img_cv, cv2.COLOR_RGB2YUV)
14     img_yuv[:, :, 0] = cv2.equalizeHist(img_yuv[:, :, 0])
15     img_norm = cv2.cvtColor(img_yuv, cv2.COLOR_YUV2RGB)
16     return Image.fromarray(img_norm)
17
18 def remove_metadata(img: Image.Image) -> Image.Image:
19     data = list(img.getdata())
20     new_img = Image.new(img.mode, img.size)
21     new_img.putdata(data)
22     return new_img
23
24 # Process raw images to clean
25 if RAW_DIR.exists():
26     print("Preprocessing raw images")
27     for filename in tqdm(sorted(os.listdir(RAW_DIR))):
28         if not filename.lower().endswith((".png", ".jpg", ".jpeg")):
29             continue
30         in_path = RAW_DIR / filename
31         out_name = Path(filename).stem + ".png"
32         out_path = CLEAN_DIR / out_name
33         try:
34             img = Image.open(in_path).convert("RGB")
35             img = center_crop(img)
36             img = img.resize((IMAGE_SIZE, IMAGE_SIZE), Image.LANCZOS)
37             img = normalize_color(img)
38             img = remove_metadata(img)
39             img.save(out_path, format="PNG")
40         except Exception as e:
41             print(f"Error processing {filename}: {e}")
42     print(f"Clean images saved to: {CLEAN_DIR}")
43
44 # Standardize damaged images
45 print("Standardizing damaged images")
46 for fname in tqdm(sorted(os.listdir(DAMAGED_DIR))):
47     p = DAMAGED_DIR / fname
48     try:
49         img = Image.open(p).convert("RGB")
50         if img.size != (IMAGE_SIZE, IMAGE_SIZE) or p.suffix.lower() != ".png":
51             img = center_crop(img).resize((IMAGE_SIZE, IMAGE_SIZE), Image.LANCZOS)
52             img = normalize_color(img)
53             img = remove_metadata(img)
54             out_name = Path(p).stem + ".png"
55             img.save(DAMAGED_DIR / out_name, format="PNG")

```

```

56         if p.suffix.lower() != ".png":
57             p.unlink(missing_ok=True)
58     except Exception as e:
59         print(f"Skipping {p}: {e}")

```

Preprocessing raw images
100%|██████████| 105/105 [00:48<00:00, 2.18it/s]
Clean images saved to: /content/drive/MyDrive/DL_Project/train/clean images
Standardizing damaged images
100%|██████████| 109/109 [00:07<00:00, 14.23it/s]

```

1 def generate_optimized_damage_mask(img_pil, debug=False):
2     """Optimized damage detection that preserves detected features"""
3     img = np.array(img_pil)
4     gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
5
6     # 1. Multi-scale Laplacian for cracks
7     crack_mask = np.zeros_like(gray, dtype=np.float32)
8     for ksize in [3, 5, 7]:
9         laplacian = cv2.Laplacian(gray, cv2.CV_64F, ksize=ksize)
10        laplacian = np.abs(laplacian)
11        crack_mask += cv2.normalize(laplacian, None, 0, 1, cv2.NORM_MINMAX)
12
13    crack_mask = cv2.normalize(crack_mask, None, 0, 255, cv2.NORM_MINMAX).astype(np.uint8)
14    _, crack_bin = cv2.threshold(crack_mask, 30, 255, cv2.THRESH_BINARY)
15
16    # 2. Sobel edges
17    sobelx = cv2.Sobel(gray, cv2.CV_64F, 1, 0, ksize=3)
18    sobely = cv2.Sobel(gray, cv2.CV_64F, 0, 1, ksize=3)
19    sobel = np.sqrt(sobelx**2 + sobely**2)
20    sobel = cv2.normalize(sobel, None, 0, 255, cv2.NORM_MINMAX).astype(np.uint8)
21    _, sobel_mask = cv2.threshold(sobel, 40, 255, cv2.THRESH_BINARY)
22
23    # 3. Canny edges
24    edges = cv2.Canny(gray, 30, 100)
25
26    # 4. LAB color space damage
27    lab = cv2.cvtColor(img, cv2.COLOR_RGB2LAB)
28    l, a, b = cv2.split(lab)
29
30    # Bright damage (paint loss)
31    _, bright_mask = cv2.threshold(l, 210, 255, cv2.THRESH_BINARY)
32
33    # Dark damage (stains, shadows from cracks)
34    _, dark_mask = cv2.threshold(l, 50, 255, cv2.THRESH_BINARY_INV)
35
36    # 5. Texture damage detection
37    blur = cv2.GaussianBlur(gray, (5, 5), 0)
38    texture_diff = cv2.absdiff(gray, blur)
39    _, texture_mask = cv2.threshold(texture_diff, 15, 255, cv2.THRESH_BINARY)
40
41    # Combine all masks with equal weight
42    combined = np.zeros_like(gray, dtype=np.float32)
43    combined += crack_bin.astype(np.float32)
44    combined += sobel_mask.astype(np.float32)
45    combined += edges.astype(np.float32)
46    combined += bright_mask.astype(np.float32)
47    combined += dark_mask.astype(np.float32)
48    combined += texture_mask.astype(np.float32)
49
50    # Normalize to 0-255
51    combined = cv2.normalize(combined, None, 0, 255, cv2.NORM_MINMAX).astype(np.uint8)
52
53    # Threshold to get binary mask
54    _, combined_binary = cv2.threshold(combined, 50, 255, cv2.THRESH_BINARY)
55
56    # MINIMAL morphological operations (preserve detail)
57    kernel_tiny = np.ones((2, 2), np.uint8)
58    kernel_small = np.ones((3, 3), np.uint8)
59
60    # Remove only the smallest noise
61    combined_binary = cv2.morphologyEx(combined_binary, cv2.MORPH_OPEN, kernel_tiny, iterations=1)
62
63    # Close small gaps
64    combined_binary = cv2.morphologyEx(combined_binary, cv2.MORPH_CLOSE, kernel_small, iterations=2)
65
66    # Moderate dilation to ensure coverage

```

```
67     combined_binary = cv2.dilate(combined_binary, kernel_small, iterations=2)
68
69     # Create soft mask with progressive blurring
70     mask_soft = combined_binary.astype(np.float32)
71     mask_soft = cv2.GaussianBlur(mask_soft, (11, 11), 0)
72     mask_soft = cv2.GaussianBlur(mask_soft, (11, 11), 0)
73
74     # Boost intensity
75     mask_soft = np.clip(mask_soft * 1.8, 0, 255).astype(np.uint8)
76
77     # Calculate coverage
78     mask_coverage = np.sum(mask_soft > 20) / mask_soft.size * 100
79
80     if debug:
81         fig, axes = plt.subplots(2, 4, figsize=(18, 9))
82         axes[0, 0].imshow(crack_bin, cmap='gray')
83         axes[0, 0].set_title('Laplacian')
84         axes[0, 0].axis('off')
85
86         axes[0, 1].imshow(sobel_mask, cmap='gray')
87         axes[0, 1].set_title('Sobel')
88         axes[0, 1].axis('off')
89
90         axes[0, 2].imshow(edges, cmap='gray')
91         axes[0, 2].set_title('Canny')
92         axes[0, 2].axis('off')
93
94         axes[0, 3].imshow(texture_mask, cmap='gray')
95         axes[0, 3].set_title('Texture')
96         axes[0, 3].axis('off')
97
98         axes[1, 0].imshow(bright_mask, cmap='gray')
99         axes[1, 0].set_title('Bright Damage')
100        axes[1, 0].axis('off')
101
102        axes[1, 1].imshow(dark_mask, cmap='gray')
103        axes[1, 1].set_title('Dark Damage')
104        axes[1, 1].axis('off')
105
106        axes[1, 2].imshow(combined_binary, cmap='gray')
107        axes[1, 2].set_title('Combined Binary')
108        axes[1, 2].axis('off')
109
110        axes[1, 3].imshow(mask_soft, cmap='gray')
111        axes[1, 3].set_title(f'Final ({mask_coverage:.1f}% coverage)')
112        axes[1, 3].axis('off')
113
114        plt.tight_layout()
115        plt.show()
116
117    print(f"Mask coverage: {mask_coverage:.2f}%")
118
119    return mask_soft
120
121 # Test on the problematic image
122 test_files = sorted(os.listdir(DAMAGED_DIR))
123 test_file = test_files[0] # or specify the one from your screenshot
124
125 print(f"Testing on: {test_file}")
126 test_img = Image.open(DAMAGED_DIR / test_file).convert("RGB")
127
128 test_mask = generate_optimized_damage_mask(test_img, debug=True)
129
130 # Show results
131 plt.figure(figsize=(18, 5))
132
133 plt.subplot(1, 4, 1)
134 plt.imshow(test_img)
135 plt.title("Damaged Image", fontsize=14)
136 plt.axis('off')
137
138 plt.subplot(1, 4, 2)
139 plt.imshow(test_mask, cmap='gray')
140 plt.title("Generated Mask", fontsize=14)
141 plt.axis('off')
142
143 plt.subplot(1, 4, 3)
```

```

144 overlay = np.array(test_img).copy()
145 mask_colored = cv2.applyColorMap(test_mask, cv2.COLORMAP_JET)
146 overlay = cv2.addWeighted(overlay, 0.7, mask_colored, 0.3, 0)
147 plt.imshow(overlay)
148 plt.title("Mask Overlay", fontsize=14)
149 plt.axis('off')
150
151 plt.subplot(1, 4, 4)
152 # Show mask intensity histogram
153 plt.hist(test_mask.flatten(), bins=50, color='blue', alpha=0.7)
154 plt.xlabel('Intensity')
155 plt.ylabel('Pixel Count')
156 plt.title('Mask Intensity Distribution')
157 plt.grid(alpha=0.3)
158
159 plt.tight_layout()
160 plt.show()
161
162 print("\n" + "*"*60)
163 print("Mask Statistics:")
164 print(f"Mean intensity: {np.mean(test_mask):.2f}")
165 print(f"Max intensity: {np.max(test_mask)}")
166 print(f"Non-zero pixels: {np.count_nonzero(test_mask)} ({(np.count_nonzero(test_mask)/test_mask.size*100:.2f)%})")
167 print("*"*60)

```

Testing on: 001.png



Mask coverage: 81.04%



```
=====
Mask Statistics:
Mean intensity: 185.01
Max intensity: 255
Non-zero pixels: 226,904 (86.56%)
=====
```

```

1 print("Generating masks for all images...")
2
3 mask_stats = []
4
5 for fname in tqdm(sorted(os.listdir(DAMAGED_DIR))):
6     if not fname.lower().endswith((".png", ".jpg", ".jpeg")):
7         continue
8
9     damaged_path = DAMAGED_DIR / fname
10    img_pil = Image.open(damaged_path).convert("RGB")
11
12    mask_np = generate_optimized_damage_mask(img_pil, debug=False)
13    mask_pil = Image.fromarray(mask_np).convert("L")
14
15    # Track statistics
16    coverage = np.sum(mask_np > 20) / mask_np.size * 100
17    mean_intensity = np.mean(mask_np)
18    mask_stats.append({
19        'file': fname,
20        'coverage': coverage,
21        'mean_intensity': mean_intensity
22    })
23
24    save_path = MASK_DIR / fname
25    mask_pil.save(save_path)
26
27 print("\nMask generation complete!")
28 print(f"\nOverall Statistics:")
29 print(f"Average coverage: {np.mean([s['coverage'] for s in mask_stats]):.2f}%")
30 print(f"Average intensity: {np.mean([s['mean_intensity'] for s in mask_stats]):.2f}")
31
32 # Show images with lowest coverage (potential problems)
33 mask_stats_sorted = sorted(mask_stats, key=lambda x: x['coverage'])
34 print(f"\nImages with lowest mask coverage (check these):")
35 for stat in mask_stats_sorted[:5]:
36     print(f" {stat['file']}: {stat['coverage']:.2f}% coverage")

```

Generating masks for all images...

Index	File	Coverage (%)
0	[REDACTED]	81.04%
1	[REDACTED]	89.05%
2	[REDACTED]	98.17%
3	[REDACTED]	73.81%
4	[REDACTED]	99.89%
5	[REDACTED]	94.20%
6	[REDACTED]	98.21%
7	[REDACTED]	87.61%
8	[REDACTED]	98.45%
9	[REDACTED]	99.96%
10	[REDACTED]	99.92%
11	[REDACTED]	99.50%
12	[REDACTED]	88.70%
13	[REDACTED]	100.00%
14	[REDACTED]	90.05%
15	[REDACTED]	87.21%
16	[REDACTED]	100.00%
17	[REDACTED]	92.55%
18	[REDACTED]	99.96%
19	[REDACTED]	88.49%
20	[REDACTED]	83.38%
21	[REDACTED]	99.68%
22	[REDACTED]	95.90%
23	[REDACTED]	96.86%
24	[REDACTED]	97.27%
25	[REDACTED]	71.76%
26	[REDACTED]	97.52%
27	[REDACTED]	79.01%
28	[REDACTED]	99.67%
29	[REDACTED]	99.72%
30	[REDACTED]	99.83%
31	[REDACTED]	78.66%
32	[REDACTED]	99.62%

```
39% |██████████| 43/109 [00:02<00:04, 13.50it/s]Mask coverage: 94.85%
Mask coverage: 79.89%
41% |██████████| 45/109 [00:02<00:05, 11.42it/s]Mask coverage: 86.08%
Mask coverage: 92.11%
Mask coverage: 96.35%
43% |██████████| 47/109 [00:02<00:06, 9.93it/s]Mask coverage: 87.87%
45% |██████████| 49/109 [00:03<00:06, 9.58it/s]Mask coverage: 95.33%
Mask coverage: 84.05%
47% |██████████| 51/109 [00:03<00:06, 9.34it/s]Mask coverage: 79.39%
48% |██████████| 52/109 [00:03<00:06, 9.10it/s]Mask coverage: 96.58%
49% |██████████| 53/109 [00:03<00:06, 9.05it/s]Mask coverage: 84.18%
50% |██████████| 54/109 [00:03<00:06, 8.79it/s]Mask coverage: 90.15%
50% |██████████| 55/109 [00:03<00:06, 8.90it/s]Mask coverage: 98.67%
Mask coverage: 99.74%
```

```
1 # Visually inspect several masks
2 num_samples = 6
3 sample_files = random.sample([f for f in os.listdir(DAMAGED_DIR) if f.endswith('.png')], min(num_samples, len(os.listdir(DA
4
5 fig, axes = plt.subplots(num_samples, 3, figsize=(12, num_samples*3))
6
7 for idx, fname in enumerate(sample_files):
8     damaged = Image.open(DAMAGED_DIR / fname).convert("RGB")
9     mask = Image.open(MASK_DIR / fname).convert("L")
10
11     # Show damaged
12     axes[idx, 0].imshow(damaged)
13     axes[idx, 0].set_title(f"{fname}", fontsize=10)
14     axes[idx, 0].axis('off')
15
16     # Show mask
17     axes[idx, 1].imshow(mask, cmap='gray')
18     coverage = np.sum(np.array(mask) > 20) / np.array(mask).size * 100
19     axes[idx, 1].set_title(f"Mask ({coverage:.1f}%)", fontsize=10)
20     axes[idx, 1].axis('off')
21
22     # Show overlay
23     overlay = np.array(damaged).copy()
24     mask_np = np.array(mask)
25     mask_colored = cv2.applyColorMap(mask_np, cv2.COLORMAP_JET)
26     overlay = cv2.addWeighted(overlay, 0.7, mask_colored, 0.3, 0)
27     axes[idx, 2].imshow(overlay)
28     axes[idx, 2].set_title("Overlay", fontsize=10)
29     axes[idx, 2].axis('off')
30
31 plt.tight_layout()
32 plt.show()
```



040.png

Mask (99.7%)



Overlay



010.png

Mask (100.0%)



Overlay



080.png

Mask (98.4%)



Overlay



075.png

Mask (97.8%)



Overlay

```

1 # setup model config
2
3 print("Loading Stable Diffusion Inpainting model...")
4
5 pipe = StableDiffusionInpaintPipeline.from_pretrained(
6     "runwayml/stable-diffusion-inpainting",
7     torch_dtype=torch.float16,
8     use_safetensors=False,
9     variant=None
10 ).to(DEVICE)
11
12 pipe.scheduler = DDPMscheduler.from_pretrained(
13     "runwayml/stable-diffusion-inpainting",
14     subfolder="scheduler"
15 )
16
17 # Optimize memory
18 pipe.safety_checker = None
19 pipe.requires_safety_checker = False
20 pipe.enable_attention_slicing()
21 pipe.enable_vae_slicing()
22
23 # Freeze VAE and text encoder
24 pipe.vae.eval()

```

```

25 pipe.text_encoder.eval()
26 pipe.vae.requires_grad_(False)
27 pipe.text_encoder.requires_grad_(False)
28
29 # UNet in FP32 for training stability
30 pipe.unet.to(device=DEVICE, dtype=torch.float32)
31 pipe.unet.requires_grad_(False)
32
33 print("Base model ready!")

```

```

Loading diffusion inpainting model...
/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
model_index.json: 100%                                         548/548 [00:00<00:00, 58.0kB/s]
Fetching 16 files: 100%                                         16/16 [01:21<00:00, 5.44s/it]
config.json:      4.78k/? [00:00<00:00, 46.6kB/s]
config.json: 100%                                         748/748 [00:00<00:00, 7.23kB/s]
config.json: 100%                                         617/617 [00:00<00:00, 6.39kB/s]
preprocessor_config.json: 100%                                         342/342 [00:00<00:00, 2.94kB/s]
merges.txt:      525k/? [00:00<00:00, 9.13MB/s]
scheduler_config.json: 100%                                         313/313 [00:00<00:00, 23.7kB/s]
safety_checker/pytorch_model.bin: 100%                                         1.22G/1.22G [00:49<00:00, 8.11MB/s]
text_encoder/pytorch_model.bin: 100%                                         492M/492M [00:35<00:00, 9.33MB/s]
tokenizer_config.json: 100%                                         806/806 [00:00<00:00, 38.6kB/s]
special_tokens_map.json: 100%                                         472/472 [00:00<00:00, 14.4kB/s]
vocab.json:      1.06M/? [00:00<00:00, 13.5MB/s]
config.json: 100%                                         552/552 [00:00<00:00, 17.2kB/s]
unet/diffusion_pytorch_model.bin: 100%                                         3.44G/3.44G [01:21<00:00, 133MB/s]
vae/diffusion_pytorch_model.bin: 100%                                         335M/335M [00:32<00:00, 13.1MB/s]
Loading pipeline components...: 100%                                         7/7 [00:27<00:00, 3.98s/it]
`torch_dtype` is deprecated! Use `dtype` instead!
Base model ready!

```



```

1 # Enhanced LoRA configuration
2 lora_config = LoraConfig(
3     r=32,    # Increased rank for better capacity
4     lora_alpha=64, # Increased alpha for stronger adaptation
5     target_modules=["to_q", "to_k", "to_v", "to_out.0"],
6     lora_dropout=0.1,
7     bias="none"
8 )
9
10 pipe.unet = get_peft_model(pipe.unet, lora_config)
11 pipe.unet.print_trainable_parameters()
12
13 print("Enhanced LoRA attached successfully!")

```

```

trainable params: 6,377,472 || all params: 865,912,836 || trainable%: 0.7365
Enhanced LoRA attached successfully!

```

```

1 # dataset and dataloader
2
3 class RestorationDataset(Dataset):
4     def __init__(self):
5         clean_files = set(os.listdir(CLEAN_DIR))
6         damaged_files = set(os.listdir(DAMAGED_DIR))
7         mask_files = set(os.listdir(MASK_DIR))
8
9         self.files = sorted(clean_files & damaged_files & mask_files)
10        print(f"Total aligned samples: {len(self.files)}")
11

```

```

1 self.img_tf = transforms.Compose([
2     transforms.Resize((512, 512)),
3     transforms.ToTensor(),
4 ])
5
6 self.mask_tf = transforms.Compose([
7     transforms.Resize((512, 512)),
8     transforms.ToTensor(),
9 ])
10
11 def __len__(self):
12     return len(self.files)
13
14 def __getitem__(self, idx):
15     fname = self.files[idx]
16
17     clean = Image.open(CLEAN_DIR / fname).convert("RGB")
18     damaged = Image.open(DAMAGED_DIR / fname).convert("RGB")
19     mask = Image.open(MASK_DIR / fname).convert("L")
20
21     return {
22         "clean": self.img_tf(clean),
23         "damaged": self.img_tf(damaged),
24         "mask": self.mask_tf(mask),
25     }
26
27
28 dataset = RestorationDataset()
29
30 loader = DataLoader(
31     dataset,
32     batch_size=1,
33     shuffle=True,
34     num_workers=0,
35     pin_memory=True
36 )
37
38 print(f"DataLoader ready | Batches: {len(loader)}")

```

Total aligned samples: 105
 DataLoader ready | Batches: 105

```

1 pipe.unet.train()
2
3 LR = 5e-6 # Lower learning rate for stability
4 EPOCHS = 10 # More epochs for better learning
5 ACCUMULATION_STEPS = 4 # Gradient accumulation
6
7 optimizer = AdamW(pipe.unet.parameters(), lr=LR, weight_decay=1e-4)
8
9 # Learning rate scheduler
10 from torch.optim.lr_scheduler import CosineAnnealingLR
11 scheduler = CosineAnnealingLR(optimizer, T_max=EPOCHS, eta_min=1e-7)
12
13 print("Starting enhanced training...")
14
15 for epoch in range(EPOCHS):
16     total_loss = 0.0
17     valid_steps = 0
18
19     progress_bar = tqdm(loader, desc=f"Epoch {epoch+1}/{EPOCHS}")
20
21     for batch_idx, batch in enumerate(progress_bar):
22         clean = batch["clean"].to(DEVICE).to(torch.float32)
23         damaged = batch["damaged"].to(DEVICE).to(torch.float32)
24         mask = torch.clamp(batch["mask"].to(DEVICE).to(torch.float32), 0, 1)
25
26         # Encode to latents
27         with torch.no_grad():
28             clean_latents = pipe.vae.encode(clean.to(torch.float16)).latent_dist.sample().to(torch.float32) * 0.18215
29             damaged_latents = pipe.vae.encode(damaged.to(torch.float16)).latent_dist.sample().to(torch.float32) * 0.18215
30
31         # Resize mask to latent space
32         mask_latent = F.interpolate(mask, size=clean_latents.shape[-2:], mode="nearest")
33         mask_latent = (mask_latent > 0.3).to(torch.float32) # Lower threshold for better coverage
34
35         # Sample noise and timesteps

```

```

36     noise = torch.randn_like(clean_latents)
37     timesteps = torch.randint(
38         0,
39         pipe.scheduler.config.num_train_timesteps,
40         (clean_latents.shape[0],),
41         device=DEVICE
42     ).long()
43
44     # Add noise to clean latents
45     noisy_latents = pipe.scheduler.add_noise(clean_latents, noise, timesteps)
46
47     # Create masked input
48     masked_latents = damaged_latents * (1 - mask_latent) + noisy_latents * mask_latent
49
50     # Concatenate inputs
51     latent_model_input = torch.cat([noisy_latents, masked_latents, mask_latent], dim=1)
52
53     # Text conditioning
54     prompts = ["professional photo restoration, repair damage, preserve details"] * clean_latents.shape[0]
55     inputs = pipe.tokenizer(
56         prompts,
57         padding="max_length",
58         truncation=True,
59         max_length=pipe.tokenizer.model_max_length,
60         return_tensors="pt"
61     )
62
63     input_ids = inputs.input_ids.to(DEVICE)
64
65     with torch.no_grad():
66         encoder_hidden_states = pipe.text_encoder(input_ids)[0].to(torch.float32)
67
68     # Forward pass
69     model_pred = pipe.unet(
70         latent_model_input,
71         timesteps,
72         encoder_hidden_states=encoder_hidden_states
73     ).sample
74
75     # Enhanced loss computation
76     # Weighted MSE emphasizing damaged regions
77     weight_mask = mask_latent * 3.0 + 0.5 # Higher weight on damaged areas
78     mse_loss = F.mse_loss(model_pred * weight_mask, noise * weight_mask)
79
80     # Perceptual loss component (L1 loss on predictions)
81     l1_loss = F.l1_loss(model_pred, noise)
82
83     # Combined loss
84     loss = mse_loss + 0.1 * l1_loss
85
86     # Gradient accumulation
87     loss = loss / ACCUMULATION_STEPS
88     loss.backward()
89
90     if (batch_idx + 1) % ACCUMULATION_STEPS == 0:
91         torch.nn.utils.clip_grad_norm_(pipe.unet.parameters(), 1.0)
92         optimizer.step()
93         optimizer.zero_grad()
94
95     total_loss += loss.item() * ACCUMULATION_STEPS
96     valid_steps += 1
97
98     progress_bar.set_postfix({"loss": loss.item() * ACCUMULATION_STEPS})
99
100    # Step scheduler
101    scheduler.step()
102
103    avg_loss = total_loss / max(1, valid_steps)
104    print(f"Epoch {epoch+1} DONE | Avg Loss: {avg_loss:.6f} | LR: {scheduler.get_last_lr()[0]:.2e}")
105
106 # Save trained model
107 SAVE_PATH = "/content/drive/MyDrive/DL_Project/Final Trained Enhanced"
108 os.makedirs(SAVE_PATH, exist_ok=True)
109
110 pipe.unet.save_pretrained(SAVE_PATH)
111 print(f"Enhanced LoRA saved at: {SAVE_PATH}")

```

```

Starting enhanced training...
Epoch 1/10: 100%|██████████| 105/105 [01:38<00:00,  1.06it/s, loss=1.55]
Epoch 1 DONE | Avg Loss: 3.146416 | LR: 4.88e-06
Epoch 2/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=0.742]
Epoch 2 DONE | Avg Loss: 2.818875 | LR: 4.53e-06
Epoch 3/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=3.77]
Epoch 3 DONE | Avg Loss: 3.068509 | LR: 3.99e-06
Epoch 4/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=2.4]
Epoch 4 DONE | Avg Loss: 2.894699 | LR: 3.31e-06
Epoch 5/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=1.82]
Epoch 5 DONE | Avg Loss: 2.713124 | LR: 2.55e-06
Epoch 6/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=4.91]
Epoch 6 DONE | Avg Loss: 3.410899 | LR: 1.79e-06
Epoch 7/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=0.134]
Epoch 7 DONE | Avg Loss: 2.518746 | LR: 1.11e-06
Epoch 8/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=0.0931]
Epoch 8 DONE | Avg Loss: 2.313404 | LR: 5.68e-07
Epoch 9/10: 100%|██████████| 105/105 [01:35<00:00,  1.10it/s, loss=1.39]
Epoch 9 DONE | Avg Loss: 2.132528 | LR: 2.20e-07
Epoch 10/10: 100%|██████████| 105/105 [01:36<00:00,  1.09it/s, loss=2.85] Epoch 10 DONE | Avg Loss: 2.583710 | LR: 1.00e-07
Enhanced LoRA saved at: /content/drive/MyDrive/DL_Project/Final Trained Enhanced

```

```

1 print("Loading model for inference")
2
3 # Load base model
4 pipe = StableDiffusionInpaintPipeline.from_pretrained(
5     "runwayml/stable-diffusion-inpainting",
6     torch_dtype=torch.float16 if DEVICE == "cuda" else torch.float32
7 ).to(DEVICE)
8
9 pipe.safety_checker = None
10 pipe.requires_safety_checker = False
11
12 # Load trained LoRA
13 pipe.load_lora_weights("/content/drive/MyDrive/DL_Project/Final Trained Enhanced")
14
15 pipe.enable_attention_slicing()
16 pipe.enable_vae_slicing()
17
18 pipe.unet.eval()
19 pipe.vae.eval()
20 pipe.text_encoder.eval()
21
22 print("Trained LoRA loaded for testing!")

```

```

Loading model for inference
Loading pipeline components...: 100%                                         7/7 [00:16<00:00,  2.73s/it]

An error occurred while trying to fetch /root/.cache/huggingface/hub/models--runwayml--stable-diffusion-inpainting/snapshots/8a4
Defaulting to unsafe serialization. Pass `allow_pickle=False` to raise an error instead.
An error occurred while trying to fetch /root/.cache/huggingface/hub/models--runwayml--stable-diffusion-inpainting/snapshots/8a4
Defaulting to unsafe serialization. Pass `allow_pickle=False` to raise an error instead.
No LoRA keys associated to UNet2DConditionModel found with the prefix='unet'. This is safe to ignore if LoRA state dict didn't o
No LoRA keys associated to CLIPTextModel found with the prefix='text_encoder'. This is safe to ignore if LoRA state dict didn't
Trained LoRA loaded for testing!

```

```

1 # Create comparison directory
2 COMPARISON_DIR = Path("/content/drive/MyDrive/DL_Project/comparisons")
3 COMPARISON_DIR.mkdir(parents=True, exist_ok=True)
4
5 print(f"Creating comparison images")
6
7 damaged_files = sorted([f for f in os.listdir(DAMAGED_DIR) if f.lower().endswith('.png', '.jpg', '.jpeg')])
8
9 for fname in tqdm(damaged_files[:20], desc="Creating comparisons"): # First 20 for quick preview
10     try:
11         # Load all versions
12         damaged_pil = Image.open(DAMAGED_DIR / fname).convert("RGB").resize((512, 512))
13         mask_pil = Image.open(MASK_DIR / fname).convert("L").resize((512, 512))
14         clean_pil = Image.open(CLEAN_DIR / fname).convert("RGB").resize((512, 512))
15         restored_name = Path(fname).stem + "_restored.png"
16         restored_pil = Image.open(RESTORED_DIR / restored_name).convert("RGB")
17
18         # Create comparison grid
19         fig, axes = plt.subplots(1, 4, figsize=(20, 5))
20

```

```

21     axes[0].imshow(damaged_pil)
22     axes[0].set_title("Damaged", fontsize=14)
23     axes[0].axis("off")
24
25     axes[1].imshow(mask_pil, cmap="gray")
26     axes[1].set_title("Mask", fontsize=14)
27     axes[1].axis("off")
28
29     axes[2].imshow(restored_pil)
30     axes[2].set_title("Restored", fontsize=14)
31     axes[2].axis("off")
32
33     axes[3].imshow(clean_pil)
34     axes[3].set_title("Original Clean", fontsize=14)
35     axes[3].axis("off")
36
37     plt.tight_layout()
38
39     # Save comparison
40     comparison_name = Path(fname).stem + "_comparison.png"
41     plt.savefig(COMPARISON_DIR / comparison_name, dpi=150, bbox_inches='tight')
42     plt.close()
43
44 except Exception as e:
45     print(f"Error creating comparison for {fname}: {e}")
46
47 print(f"\nComparison images saved to: {COMPARISON_DIR}")

```

Creating comparison images
 Creating comparisons: 0% | 0/20 [00:00<00:, ?it/s] Error creating comparison for 001.png: name 'RESTORED_DIR' is not defined
 Creating comparisons: 10% | 2/20 [00:00<00:02, 8.01it/s] Error creating comparison for 002.png: [Errno 2] No such file or directory: '/content/drive/MyDrive/DL_Project/train/clean image 002.png'
 Creating comparisons: 30% | 6/20 [00:00<00:00, 18.09it/s] Error creating comparison for 003.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 004.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 005.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 006.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 007.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 009.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 010.png: name 'RESTORED_DIR' is not defined
 Creating comparisons: 70% | 14/20 [00:00<00:00, 26.73it/s] Error creating comparison for 011.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 012.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 013.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 014.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 015.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 016.png: name 'RESTORED_DIR' is not defined
 Creating comparisons: 100% | 20/20 [00:00<00:00, 24.51it/s] Error creating comparison for 017.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 018.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 019.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 020.png: name 'RESTORED_DIR' is not defined
 Error creating comparison for 021.png: [Errno 2] No such file or directory: '/content/drive/MyDrive/DL_Project/train/clean image 021.png'

Comparison images saved to: /content/drive/MyDrive/DL_Project/comparisons

```

1 def hybrid_restoration(damaged_pil, mask_pil, use_ai=True):
2     """
3         Combines OpenCV inpainting (for style preservation)
4         with light AI enhancement (for quality)
5     """
6     damaged_np = np.array(damaged_pil).astype(np.float32)
7     mask_np = np.array(mask_pil).astype(np.float32) / 255.0
8
9     # Check if mask has enough content
10    mask_coverage = np.sum(mask_np > 0.1) / mask_np.size * 100
11
12    if mask_coverage < 1.0:
13        # Too little damage, return original
14        return damaged_pil
15
16    # Step 1: OpenCV inpainting (style-preserving base)
17    damaged_uint8 = damaged_np.astype(np.uint8)
18    mask_uint8 = (mask_np * 255).astype(np.uint8)
19    mask_inv = 255 - mask_uint8
20
21    inpainted_ns = cv2.inpaint(damaged_uint8, mask_inv, 7, cv2.INPAINT_NS)
22    inpainted_teela = cv2.inpaint(damaged_uint8, mask_inv, 7, cv2.INPAINT_TELEA)
23
24    # Blend both OpenCV methods
25    cv_result = (0.6 * inpainted_ns + 0.4 * inpainted_teela).astype(np.uint8)

```

```

26
27     if use_ai and mask_coverage > 5.0:
28         # Step 2: Light AI refinement only if significant damage
29         cv_result_pil = Image.fromarray(cv_result)
30
31         # Reduce mask strength for AI
32         mask_reduced = (mask_np * 0.5).clip(0, 1)
33         mask_reduced = cv2.GaussianBlur(mask_reduced, (21, 21), 0)
34         mask_reduced_pil = Image.fromarray((mask_reduced * 255).astype(np.uint8)).convert("L")
35
36         with torch.no_grad():
37             ai_result = pipe(
38                 prompt="subtle texture repair, match surrounding style exactly, painting texture",
39                 negative_prompt="photorealistic, photograph, smooth, modern, different face, identity change",
40                 image=cv_result_pil,
41                 mask_image=mask_reduced_pil,
42                 num_inference_steps=30,
43                 guidance_scale=3.5,
44                 strength=0.45
45             ).images[0]
46
47         ai_np = np.array(ai_result).astype(np.float32)
48         cv_np = np.array(cv_result).astype(np.float32)
49
50         # Blend AI with OpenCV (60% OpenCV, 40% AI)
51         blended = cv_np * 0.6 + ai_np * 0.4
52     else:
53         blended = cv_result.astype(np.float32)
54
55     # Step 3: Final blending with original
56     mask_smooth = cv2.GaussianBlur(mask_np, (31, 31), 0)
57     mask_smooth = np.clip(mask_smooth * 0.8, 0, 1) # Reduce to 80%
58
59     final = blended * mask_smooth[..., None] + damaged_np * (1 - mask_smooth[..., None])
60
61     # Color correction to match original
62     for c in range(3):
63         orig_mean = np.mean(damaged_np[:, :, c])
64         orig_std = np.std(damaged_np[:, :, c])
65         final_mean = np.mean(final[:, :, c])
66         final_std = np.std(final[:, :, c])
67
68         if final_std > 0:
69             final[:, :, c] = (final[:, :, c] - final_mean) * (orig_std / final_std) + orig_mean
70
71     final = np.clip(final, 0, 255).astype(np.uint8)
72
73     return Image.fromarray(final)
74
75 # Test
76 fname = "005.png"
77 print(f"Testing hybrid restoration on: {fname}")
78
79 damaged_pil = Image.open(DAMAGED_DIR / fname).convert("RGB").resize((512, 512))
80 mask_pil = Image.open(MASK_DIR / fname).convert("L").resize((512, 512))
81 clean_pil = Image.open(CLEAN_DIR / fname).convert("RGB").resize((512, 512))
82
83 # Try both modes
84 restored_cv_only = hybrid_restoration(damaged_pil, mask_pil, use_ai=False)
85 restored_hybrid = hybrid_restoration(damaged_pil, mask_pil, use_ai=True)
86
87 # Metrics
88 clean_np = np.array(clean_pil).astype(np.float32)
89 damaged_np = np.array(damaged_pil).astype(np.float32)
90 cv_np = np.array(restored_cv_only).astype(np.float32)
91 hybrid_np = np.array(restored_hybrid).astype(np.float32)
92
93 psnr_d = psnr(clean_np, damaged_np, data_range=255)
94 psnr_cv = psnr(clean_np, cv_np, data_range=255)
95 psnr_h = psnr(clean_np, hybrid_np, data_range=255)
96
97 ssim_d = ssim(clean_np, damaged_np, data_range=255, channel_axis=2)
98 ssim_cv = ssim(clean_np, cv_np, data_range=255, channel_axis=2)
99 ssim_h = ssim(clean_np, hybrid_np, data_range=255, channel_axis=2)
100
101 # Display
102 fig, axes = plt.subplots(1, 5, figsize=(22, 5))

```

```

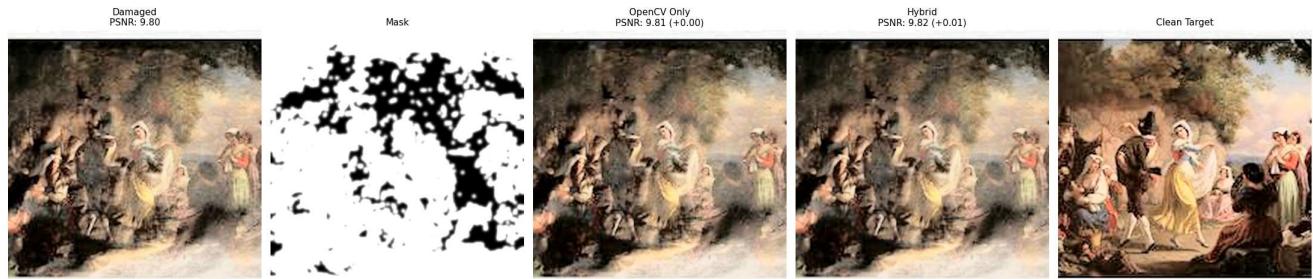
103
104 axes[0].imshow(damaged_pil)
105 axes[0].set_title(f"Damaged\nPSNR: {psnr_d:.2f}", fontsize=11)
106 axes[0].axis('off')
107
108 axes[1].imshow(mask_pil, cmap='gray')
109 axes[1].set_title("Mask", fontsize=11)
110 axes[1].axis('off')
111
112 axes[2].imshow(restored_cv_only)
113 axes[2].set_title(f"OpenCV Only\nPSNR: {psnr_cv:.2f} ({psnr_cv-psnr_d:+.2f})", fontsize=11)
114 axes[2].axis('off')
115
116 axes[3].imshow(restored_hybrid)
117 axes[3].set_title(f"Hybrid\nPSNR: {psnr_h:.2f} ({psnr_h-psnr_d:+.2f})", fontsize=11)
118 axes[3].axis('off')
119
120 axes[4].imshow(clean_pil)
121 axes[4].set_title("Clean Target", fontsize=11)
122 axes[4].axis('off')
123
124 plt.tight_layout()
125 plt.show()
126
127 print(f"\nOpenCV Only: PSNR {psnr_cv-psnr_d:+.2f} dB, SSIM {ssim_cv-ssim_d:+.4f}")
128 print(f"Hybrid:      PSNR {psnr_h-psnr_d:+.2f} dB, SSIM {ssim_h-ssim_d:+.4f}")

```

Testing hybrid restoration on: 005.png

100%

13/13 [00:02<00:00, 4.84it/s]



OpenCV Only: PSNR +0.00 dB, SSIM +0.0083
 Hybrid: PSNR +0.01 dB, SSIM +0.0110

```

1 RESTORED_DIR = Path("/content/drive/MyDrive/DL_Project/restored_images_hybrid")
2 RESTORED_DIR.mkdir(parents=True, exist_ok=True)
3
4 print("Generating hybrid restorations...")
5
6 use_ai = True # Set to False for OpenCV-only (faster)
7
8 for fname in tqdm(sorted(os.listdir(DAMAGED_DIR))):
9     if not fname.lower().endswith('.png', '.jpg', '.jpeg'):
10         continue
11
12     try:
13         damaged_pil = Image.open(DAMAGED_DIR / fname).convert("RGB").resize((512, 512))
14         mask_pil = Image.open(MASK_DIR / fname).convert("L").resize((512, 512))
15
16         restored = hybrid_restoration(damaged_pil, mask_pil, use_ai=use_ai)
17
18         output_name = Path(fname).stem + "_restored.png"
19         restored.save(RESTORED_DIR / output_name)
20
21     except Exception as e:
22         print(f"\nError: {fname}: {e}")
23
24 print(f"\nSaved to: {RESTORED_DIR}")

```



```

Generating hybrid restorations...
0% | 0/109 [00:00<?, ?it/s]
100%                               13/13 [00:02<00:00,  5.06it/s]
1% | 1/109 [00:04<07:55,  4.40s/it]
100%                               13/13 [00:02<00:00,  5.03it/s]
2% | 2/109 [00:07<06:38,  3.72s/it]
100%                               13/13 [00:02<00:00,  4.97it/s]
3% | 3/109 [00:12<07:09,  4.05s/it]
100%                               13/13 [00:02<00:00,  4.98it/s]
4% | 4/109 [00:16<07:05,  4.05s/it]
100%                               13/13 [00:02<00:00,  4.92it/s]
5% | 5/109 [00:20<07:04,  4.08s/it]
100%                               13/13 [00:02<00:00,  4.92it/s]
6% | 6/109 [00:23<06:32,  3.81s/it]
100%                               13/13 [00:02<00:00,  4.91it/s]
6% | 7/109 [00:28<06:53,  4.06s/it]
100%                               13/13 [00:02<00:00,  4.87it/s]
7% | 8/109 [00:31<06:42,  3.99s/it]
100%                               13/13 [00:02<00:00,  4.84it/s]
8% | 9/109 [00:35<06:18,  3.78s/it]
100%                               13/13 [00:02<00:00,  4.84it/s]
9% | 10/109 [00:39<06:41,  4.05s/it]
100%                               13/13 [00:02<00:00,  4.89it/s]
10% | 11/109 [00:43<06:20,  3.89s/it]
100%                               13/13 [00:02<00:00,  4.94it/s]
11% | 12/109 [00:46<05:59,  3.70s/it]
100%                               13/13 [00:02<00:00,  4.93it/s]
12% | 13/109 [00:50<05:44,  3.59s/it]
100%                               13/13 [00:02<00:00,  4.95it/s]
13% | 14/109 [00:53<05:42,  3.61s/it]
100%                               13/13 [00:02<00:00,  4.95it/s]
14% | 15/109 [00:57<05:46,  3.68s/it]
100%                               13/13 [00:02<00:00,  4.99it/s]
15% | 16/109 [01:01<05:42,  3.69s/it]
100%                               13/13 [00:02<00:00,  5.00it/s]
16% | 17/109 [01:05<05:53,  3.85s/it]
100%                               13/13 [00:02<00:00,  5.03it/s]
17% | 18/109 [01:09<05:45,  3.79s/it]
100%                               13/13 [00:02<00:00,  5.02it/s]
17% | 19/109 [01:12<05:26,  3.63s/it]
100%                               13/13 [00:02<00:00,  5.03it/s]
18% | 20/109 [01:15<05:17,  3.57s/it]
100%                               13/13 [00:02<00:00,  5.03it/s]
19% | 21/109 [01:19<05:15,  3.59s/it]
100%                               13/13 [00:02<00:00,  5.03it/s]
20% | 22/109 [01:23<05:29,  3.79s/it]
100%                               13/13 [00:02<00:00,  5.02it/s]
21% | 23/109 [01:26<05:11,  3.62s/it]
100%                               13/13 [00:02<00:00,  4.98it/s]
22% | 24/109 [01:31<05:22,  3.80s/it]
100%                               13/13 [00:02<00:00,  5.00it/s]
23% | 25/109 [01:35<05:36,  4.00s/it]
100%                               13/13 [00:02<00:00,  4.97it/s]
24% | 26/109 [01:38<05:13,  3.77s/it]
100%                               13/13 [00:02<00:00,  4.93it/s]
25% | 27/109 [01:42<05:13,  3.82s/it]
100%                               13/13 [00:02<00:00,  4.93it/s]
26% | 28/109 [01:46<04:59,  3.70s/it]
100%                               13/13 [00:02<00:00,  4.93it/s]
27% | 29/109 [01:50<05:10,  3.89s/it]

```

100%		13/13 [00:02<00:00, 4.95it/s]
28% ██████	30/109 [01:54<05:15, 3.99s/it]	13/13 [00:02<00:00, 4.94it/s]
100%		13/13 [00:02<00:00, 4.94it/s]
28% ██████	31/109 [01:58<05:00, 3.85s/it]	13/13 [00:02<00:00, 4.94it/s]
100%		13/13 [00:02<00:00, 4.94it/s]
29% ██████	32/109 [02:02<04:55, 3.83s/it]	13/13 [00:02<00:00, 4.97it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
30% ██████	33/109 [02:06<04:54, 3.88s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.99it/s]
31% ██████	34/109 [02:09<04:46, 3.83s/it]	13/13 [00:02<00:00, 4.99it/s]
100%		13/13 [00:02<00:00, 5.00it/s]
32% ██████	35/109 [02:14<05:12, 4.22s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
33% ██████	36/109 [02:18<04:55, 4.05s/it]	13/13 [00:02<00:00, 5.00it/s]
100%		13/13 [00:02<00:00, 5.00it/s]
34% ██████	37/109 [02:22<04:58, 4.14s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
35% ██████	38/109 [02:26<04:38, 3.92s/it]	13/13 [00:02<00:00, 4.94it/s]
100%		13/13 [00:02<00:00, 4.94it/s]
36% ██████	39/109 [02:29<04:23, 3.77s/it]	13/13 [00:02<00:00, 4.97it/s]
100%		13/13 [00:02<00:00, 4.97it/s]
37% ██████	40/109 [02:33<04:10, 3.63s/it]	13/13 [00:02<00:00, 4.96it/s]
100%		13/13 [00:02<00:00, 4.96it/s]
38% ██████	41/109 [02:37<04:23, 3.87s/it]	13/13 [00:02<00:00, 4.97it/s]
100%		13/13 [00:02<00:00, 4.97it/s]
39% ██████	42/109 [02:40<04:10, 3.75s/it]	13/13 [00:02<00:00, 4.97it/s]
100%		13/13 [00:02<00:00, 4.97it/s]
39% ██████	43/109 [02:44<04:11, 3.81s/it]	13/13 [00:02<00:00, 4.96it/s]
100%		13/13 [00:02<00:00, 4.96it/s]
40% ██████	44/109 [02:49<04:20, 4.00s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
41% ██████	45/109 [02:54<04:28, 4.20s/it]	13/13 [00:02<00:00, 4.99it/s]
100%		13/13 [00:02<00:00, 4.99it/s]
42% ██████	46/109 [02:57<04:16, 4.08s/it]	13/13 [00:02<00:00, 4.97it/s]
100%		13/13 [00:02<00:00, 4.97it/s]
43% ██████	47/109 [03:01<04:06, 3.97s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
44% ██████	48/109 [03:06<04:16, 4.20s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
45% ██████	49/109 [03:10<04:04, 4.08s/it]	13/13 [00:02<00:00, 4.96it/s]
100%		13/13 [00:02<00:00, 4.96it/s]
46% ██████	50/109 [03:14<04:01, 4.09s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
47% ██████	51/109 [03:19<04:14, 4.39s/it]	13/13 [00:02<00:00, 4.96it/s]
100%		13/13 [00:02<00:00, 4.96it/s]
48% ██████	52/109 [03:22<03:56, 4.15s/it]	13/13 [00:02<00:00, 4.98it/s]
100%		13/13 [00:02<00:00, 4.98it/s]
49% ██████	53/109 [03:27<03:53, 4.16s/it]	13/13 [00:02<00:00, 4.97it/s]
100%		13/13 [00:02<00:00, 4.97it/s]
50% ██████	54/109 [03:31<03:53, 4.24s/it]	13/13 [00:02<00:00, 4.95it/s]
100%		13/13 [00:02<00:00, 4.95it/s]
50% ██████	55/109 [03:35<03:37, 4.04s/it]	13/13 [00:02<00:00, 4.94it/s]
100%		13/13 [00:02<00:00, 4.93it/s]
51% ██████	56/109 [03:38<03:23, 3.84s/it]	13/13 [00:02<00:00, 4.93it/s]
100%		13/13 [00:02<00:00, 4.93it/s]
52% ██████	57/109 [03:41<03:13, 3.72s/it]	13/13 [00:02<00:00, 4.95it/s]
100%		13/13 [00:03<00:00, 4.54it/s]
53% ██████	58/109 [03:46<03:22, 3.96s/it]	13/13 [00:03<00:00, 4.54it/s]
100%		13/13 [00:03<00:00, 4.54it/s]

54% | [REDACTED] | 59/109 [03:50<03:21, 4.03s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
55% | [REDACTED] | 60/109 [03:55<03:23, 4.15s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
56% | [REDACTED] | 61/109 [03:59<03:19, 4.15s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
57% | [REDACTED] | 62/109 [04:02<03:03, 3.90s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
58% | [REDACTED] | 63/109 [04:06<02:56, 3.84s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.93it/s]
59% | [REDACTED] | 64/109 [04:09<02:50, 3.78s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
60% | [REDACTED] | 65/109 [04:13<02:39, 3.63s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.95it/s]
61% | [REDACTED] | 66/109 [04:16<02:34, 3.59s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
61% | [REDACTED] | 67/109 [04:21<02:45, 3.94s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.99it/s]
62% | [REDACTED] | 68/109 [04:26<02:59, 4.39s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 5.02it/s]
63% | [REDACTED] | 69/109 [04:30<02:42, 4.06s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.99it/s]
64% | [REDACTED] | 70/109 [04:33<02:36, 4.01s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 5.01it/s]
65% | [REDACTED] | 71/109 [04:39<02:45, 4.35s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.99it/s]
66% | [REDACTED] | 72/109 [04:43<02:41, 4.37s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 5.00it/s]
67% | [REDACTED] | 73/109 [04:47<02:29, 4.15s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
68% | [REDACTED] | 74/109 [04:50<02:19, 3.99s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.94it/s]
69% | [REDACTED] | 75/109 [04:54<02:11, 3.87s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.98it/s]
70% | [REDACTED] | 76/109 [04:58<02:14, 4.08s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
71% | [REDACTED] | 77/109 [05:02<02:04, 3.91s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.95it/s]
72% | [REDACTED] | 78/109 [05:06<02:04, 4.00s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
72% | [REDACTED] | 79/109 [05:10<01:56, 3.88s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
73% | [REDACTED] | 80/109 [05:14<01:52, 3.87s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
74% | [REDACTED] | 81/109 [05:18<01:53, 4.07s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.99it/s]
75% | [REDACTED] | 82/109 [05:22<01:48, 4.02s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.98it/s]
76% | [REDACTED] | 83/109 [05:26<01:40, 3.88s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
77% | [REDACTED] | 84/109 [05:31<01:44, 4.20s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.96it/s]
78% | [REDACTED] | 85/109 [05:35<01:41, 4.21s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.97it/s]
79% | [REDACTED] | 86/109 [05:38<01:31, 3.97s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.98it/s]
80% | [REDACTED] | 87/109 [05:42<01:23, 3.81s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.99it/s]
81% | [REDACTED] | 88/109 [05:46<01:23, 3.98s/it]
100% | [REDACTED] | 13/13 [00:02<00:00, 4.99it/s]

```

1 import csv
2
3 # Create metrics file
4 METRICS_PATH = Path("/content/drive/MyDrive/DL_Project/restoration_metrics.csv")
5
6 print("Computing and saving metrics for all restored images...")
7
8 # First, check what files exist
9 print("\nChecking directories...")
10 print(f"Damaged images: {len(list(DAMAGED_DIR.glob('*.*')))} files")
11 print(f"Clean images: {len(list(CLEAN_DIR.glob('*.*')))} files")
12 print(f"Restored images: {len(list(RESTORED_DIR.glob('*.*')))} files")
13
14 # Get list of restored files (they have _restored suffix)
15 restored_files = list(RESTORED_DIR.glob("*_restored.*"))
16 print(f"\nFound {len(restored_files)} restored images")
17
18 if len(restored_files) == 0:
19     print("\nERROR: No restored images found!")
20     print("Please run the restoration cell (Cell 9A) first to generate restored images.")
21     print(f"Expected location: {RESTORED_DIR}")
22 else:
23     results = []
24
25     for restored_path in tqdm(restored_files, desc="Computing metrics"):
26         try:
27             # Extract original filename (remove _restored suffix)
28             original_fname = restored_path.stem.replace("_restored", "") + ".png"
29
30             # Check if corresponding files exist
31             clean_path = CLEAN_DIR / original_fname
32             damaged_path = DAMAGED_DIR / original_fname
33
34             if not clean_path.exists():
35                 print(f"Warning: Clean image not found for {original_fname}")
36                 continue
37
38             if not damaged_path.exists():
39                 print(f"Warning: Damaged image not found for {original_fname}")
40                 continue
41
42             # Load images
43             clean_pil = Image.open(clean_path).convert("RGB").resize((512, 512))
44             damaged_pil = Image.open(damaged_path).convert("RGB").resize((512, 512))
45             restored_pil = Image.open(restored_path).convert("RGB").resize((512, 512))
46
47             # Convert to numpy
48             clean_np = np.array(clean_pil).astype(np.float32)
49             damaged_np = np.array(damaged_pil).astype(np.float32)
50             restored_np = np.array(restored_pil).astype(np.float32)
51
52             # Compute metrics
53             psnr_damaged = psnr(clean_np, damaged_np, data_range=255)
54             ssim_damaged = ssim(clean_np, damaged_np, data_range=255, channel_axis=2)
55             psnr_restored = psnr(clean_np, restored_np, data_range=255)
56             ssim_restored = ssim(clean_np, restored_np, data_range=255, channel_axis=2)
57
58             results.append({
59                 'filename': original_fname,
60                 'psnr_damaged': psnr_damaged,
61                 'ssim_damaged': ssim_damaged,
62                 'psnr_restored': psnr_restored,
63                 'ssim_restored': ssim_restored,
64                 'psnr_improvement': psnr_restored - psnr_damaged,
65                 'ssim_improvement': ssim_restored - ssim_damaged
66             })
67
68         except Exception as e:
69             print(f"Error computing metrics for {restored_path.name}: {e}")
70
71     if len(results) > 0:
72         # Save to CSV
73         with open(METRICS_PATH, 'w', newline='') as csvfile:
74             fieldnames = ['filename', 'psnr_damaged', 'ssim_damaged', 'psnr_restored',
75                         'ssim_restored', 'psnr_improvement', 'ssim_improvement']
76             writer = csv.DictWriter(csvfile, fieldnames=fieldnames)

```

```

77         writer.writeheader()
78         writer.writerows(results)
79
80     print(f"\nMetrics saved to: {METRICS_PATH}")
81
82     # Print summary statistics
83     psnr_improvements = [r['psnr_improvement'] for r in results]
84     ssim_improvements = [r['ssim_improvement'] for r in results]
85
86     print("\n" + "="*60)
87     print("SUMMARY STATISTICS")
88     print("="*60)
89     print(f"Total images processed: {len(results)}")
90
91     # Damaged vs Clean
92     avg_psnr_damaged = np.mean([r['psnr_damaged'] for r in results])
93     avg_ssim_damaged = np.mean([r['ssim_damaged'] for r in results])
94
95     # Restored vs Clean
96     avg_psnr_restored = np.mean([r['psnr_restored'] for r in results])
97     avg_ssim_restored = np.mean([r['ssim_restored'] for r in results])
98
99
100    print(f"\nDamaged → Clean | PSNR: {avg_psnr_damaged:.2f} ± {np.std([r['psnr_damaged']] for r in results):.2f} dB")
101   print(f"Restored → Clean | PSNR: {avg_psnr_restored:.2f} ± {np.std([r['psnr_restored']] for r in results):.2f} dB")
102
103   print(f"\nIMPROVEMENT:")
104   print(f"PSNR Gain: {np.mean(psnr_improvements):+.2f} ± {np.std(psnr_improvements):.2f} dB")
105   print(f"SSIM Gain: {np.mean(ssim_improvements):+.4f} ± {np.std(ssim_improvements):.4f}")
106
107   if len(psnr_improvements) > 0:
108       print(f"\nBest PSNR improvement: {max(psnr_improvements):+.2f} dB ({results[psnr_improvements.index(max(psnr_improvements))]['image']:s} | PSNR: {results[psnr_improvements.index(max(psnr_improvements))]['psnr_improvement']:.2f} dB)")
109       print(f"Worst PSNR improvement: {min(psnr_improvements):+.2f} dB ({results[psnr_improvements.index(min(psnr_improvements))]['image']:s} | PSNR: {results[psnr_improvements.index(min(psnr_improvements))]['psnr_improvement']:.2f} dB)")
110       print(f"\nBest SSIM improvement: {max(ssim_improvements):+.4f} ({results:ssim_improvements.index(max(ssim_improvements))['image']:s} | SSIM: {results:ssim_improvements.index(max(ssim_improvements))['ssim_improvement']:.4f})")
111       print(f"Worst SSIM improvement: {min(ssim_improvements):+.4f} ({results:ssim_improvements.index(min(ssim_improvements))['image']:s} | SSIM: {results:ssim_improvements.index(min(ssim_improvements))['ssim_improvement']:.4f})")
112
113   print("="*60)
114
115   # Create visualization of improvements
116   plt.figure(figsize=(14, 5))
117
118   plt.subplot(1, 2, 1)
119   plt.scatter(range(len(psnr_improvements)), psnr_improvements, alpha=0.6, s=50)
120   plt.axhline(y=0, color='r', linestyle='--', linewidth=1, label='No improvement')
121   plt.xlabel('Image Index')
122   plt.ylabel('PSNR Improvement (dB)')
123   plt.title('PSNR Improvement per Image')
124   plt.grid(alpha=0.3)
125   plt.legend()
126
127   plt.subplot(1, 2, 2)
128   plt.scatter(range(len(ssim_improvements)), ssim_improvements, alpha=0.6, color='green', s=50)
129   plt.axhline(y=0, color='r', linestyle='--', linewidth=1, label='No improvement')
130   plt.xlabel('Image Index')
131   plt.ylabel('SSIM Improvement')
132   plt.title('SSIM Improvement per Image')
133   plt.grid(alpha=0.3)
134   plt.legend()
135
136   plt.tight_layout()
137   plt.show()
138
139 else:
140     print("\nERROR: No results were computed. Check if files exist and match.")

```

```
Computing and saving metrics for all restored images...
```

```
Checking directories...
```

```
Damaged images: 109 files
```

```
Clean images: 105 files
```

```
Restored images: 109 files
```

```
Found 109 restored images
```

```
Computing metrics: 1% | 1/109 [00:00<00:23, 4.68it/s]Warning: Clean image not found for 002.png
Computing metrics: 19% | 21/109 [00:03<00:10, 8.04it/s]Warning: Clean image not found for 021.png
Computing metrics: 29% | 32/109 [00:05<00:09, 8.20it/s]Warning: Clean image not found for 032.png
Computing metrics: 72% | 78/109 [00:12<00:06, 4.71it/s]Warning: Clean image not found for 081.png
Computing metrics: 100% | 109/109 [00:18<00:00, 5.99it/s]
```

```
Metrics saved to: /content/drive/MyDrive/DL_Project/restoration_metrics.csv
```

```
=====
SUMMARY STATISTICS
=====
```

```
Total images processed: 105
```

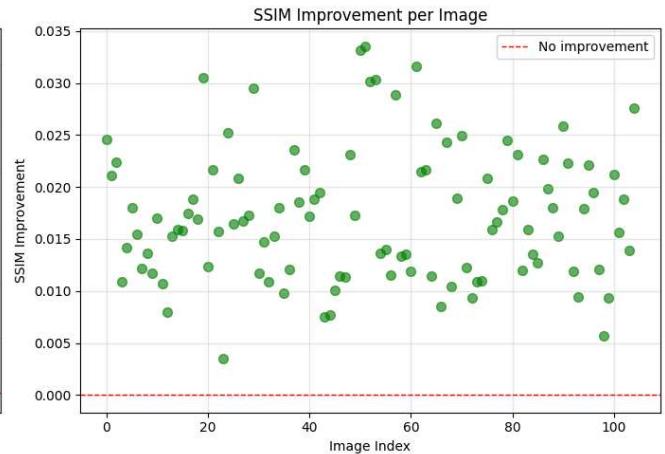
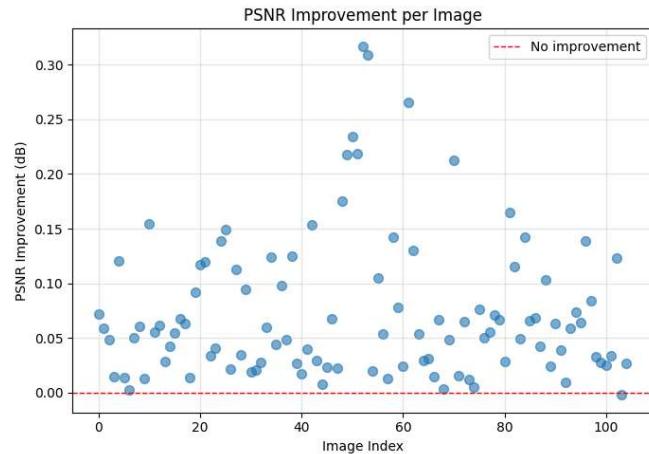
```
Damaged → Clean | PSNR: 12.50 ± 2.35 dB | SSIM: 0.2426 ± 0.1045
Restored → Clean | PSNR: 12.57 ± 2.39 dB | SSIM: 0.2598 ± 0.1063
```

IMPROVEMENT:

```
PSNR Gain: +0.07 ± 0.06 dB
SSIM Gain: +0.0172 ± 0.0063
```

```
Best PSNR improvement: +0.32 dB (058.png)
Worst PSNR improvement: -0.00 dB (110.png)
```

```
Best SSIM improvement: +0.0335 (056.png)
Worst SSIM improvement: +0.0035 (027.png)
=====
```



```
1 # Select a random test image to test restoration
2 fname = random.choice(os.listdir(DAMAGED_DIR))
3 print(f"Testing on: {fname}")
4
5 damaged_pil = Image.open(DAMAGED_DIR / fname).convert("RGB").resize((512, 512))
6 mask_pil = Image.open(MASK_DIR / fname).convert("L").resize((512, 512))
7 clean_pil = Image.open(CLEAN_DIR / fname).convert("RGB").resize((512, 512))
8
9 # Perform restoration
10 with torch.no_grad():
11     restored_pil = pipe(
12         prompt="professional museum-quality photo restoration, repair all damage, preserve original details and colors, nat",
13         negative_prompt="blur, artifacts, distortion, oversaturated, fake, synthetic, unnatural, scratches, cracks",
14         image=damaged_pil,
15         mask_image=mask_pil,
16         num_inference_steps=75,
17         guidance_scale=7.5,
18         strength=0.95
19     ).images[0]
20
21 # Display results
22 plt.figure(figsize=(20, 5))
23
24 plt.subplot(1, 4, 1)
```

```

25 plt.imshow(damaged_pil)
26 plt.title("Damaged", fontsize=14)
27 plt.axis("off")
28
29 plt.subplot(1, 4, 2)
30 plt.imshow(mask_pil, cmap="gray")
31 plt.title("Mask", fontsize=14)
32 plt.axis("off")
33
34 plt.subplot(1, 4, 3)
35 plt.imshow(restored_pil)
36 plt.title("Restored (Enhanced LoRA)", fontsize=14)
37 plt.axis("off")
38
39 plt.subplot(1, 4, 4)
40 plt.imshow(clean_pil)
41 plt.title("Original Clean", fontsize=14)
42 plt.axis("off")
43
44 plt.tight_layout()
45 plt.show()
46
47 # Compute metrics
48 p_damaged, s_damaged = psnr(np.array(clean_pil), np.array(damaged_pil), data_range=255), \
49                 ssim(np.array(clean_pil), np.array(damaged_pil), data_range=255, channel_axis=2)
50 p_restored, s_restored = psnr(np.array(clean_pil), np.array(restored_pil), data_range=255), \
51                 ssim(np.array(clean_pil), np.array(restored_pil), data_range=255, channel_axis=2)
52
53 print(f"\nDamaged → Clean | PSNR: {p_damaged:.2f} dB | SSIM: {s_damaged:.4f}")
54 print(f"Restored → Clean | PSNR: {p_restored:.2f} dB | SSIM: {s_restored:.4f}")
55 print(f"Improvement      | PSNR: {p_restored - p_damaged:+.2f} dB | SSIM: {s_restored - s_damaged:+.4f}")

```

Testing on: 086.png

100%

71/71 [00:14<00:00, 4.93it/s]



Damaged → Clean | PSNR: 16.94 dB | SSIM: 0.4425
 Restored → Clean | PSNR: 16.73 dB | SSIM: 0.5244
 Improvement | PSNR: -0.20 dB | SSIM: +0.0820

```

1 NUM_SAMPLES = 50
2
3 psnr_damaged_list = []
4 ssim_damaged_list = []
5 psnr_restored_list = []
6 ssim_restored_list = []
7
8 def tensor_to_pil(t, is_mask=False):
9     """Convert tensor to PIL image"""
10    if t.min() < 0:
11        t = (t + 1) / 2
12
13    t = t.squeeze().cpu().numpy()
14
15    # Handle grayscale mask (2D)
16    if is_mask or t.ndim == 2:
17        t = (t * 255).clip(0, 255).astype(np.uint8)
18        return Image.fromarray(t, mode='L')
19
20    # Handle RGB image (3D)
21    if t.ndim == 3:
22        # If channels first (C, H, W), convert to (H, W, C)

```